

HR Analytics Project- Understanding the Attrition in HR

Every year a lot of companies hire a number of employees. The companies invest time and money in training those employees, not just this but there are training programs within the companies for their existing employees as well. The aim of these programs is to increase the effectiveness of their employees. But where HR Analytics fit in this? and is it just about improving the performance of employees?

Attrition affecting Companies

A major problem in high employee attrition is its cost to an organization. Job postings, hiring processes, paperwork, and new hire training are some of the common expenses of losing employees and replacing them. Additionally, regular employee turnover prohibits your organization from increasing its collective knowledge base and experience over time. This is especially concerning if your business is customer-facing, as customers often prefer to interact with familiar people. Errors and issues are more likely if you constantly have new workers.

```
# Importing packages
```

```
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
```

```
#reading a dataset
```

```
df = pd.read_csv('WA_Fn-UseC_-HR-Employee-Attrition.csv')
df.head()
```

	Age	Attrition	...	YearsSinceLastPromotion	YearsWithCurrManager
0	41	Yes	...	0	5
1	49	No	...	1	7
2	37	Yes	...	0	0
3	33	No	...	3	0
4	27	No	...	2	2

```
[5 rows x 35 columns]
```

with the above observation of dataset,our target variable is Attrition and we are going to predict Attrition considering the features like Department,EducationField ,etc., which are given in the dataset

```
df.duplicated().sum()
```

```
0
```

So,there are no duplicate records

```
df['Over18'].value_counts()
```

```
Y      1470
```

```
Name: Over18, dtype: int64
```

```
#with that,we can remove column "Over18" which is consisting of just one value as "y"
```

```
df = df.drop(['Over18'],axis=1)
```

```
df['StandardHours'].value_counts()
```

```
80      1470
```

```
Name: StandardHours, dtype: int64
```

```
#with that,we can remove column "StandardHours" which is consisting of just one value as "80"
```

```
df = df.drop(['StandardHours'],axis=1)
```

```
df['EmployeeCount'].value_counts()
```

```
1      1470
```

```
Name: EmployeeCount, dtype: int64
```

```
#with that,we can remove column "EmployeeCount" which is consisting of just one value as "1"
```

```
df = df.drop(['EmployeeCount'],axis=1)
```

```
#we can also drop "EmployeeNumber" column which is consisting of unique id numbers for each employee.
```

```
df = df.drop(['EmployeeNumber'],axis=1)
```

```
#to check how many rows and columns in dataframe
```

```
df.shape
```

```
(1470, 31)
```

```
#to get the info of this dataframe
```

```
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
```

```
RangeIndex: 1470 entries, 0 to 1469
```

```
Data columns (total 31 columns):
```

#	Column	Non-Null Count	Dtype
0	Age	1470 non-null	int64
1	Attrition	1470 non-null	object
2	BusinessTravel	1470 non-null	object
3	DailyRate	1470 non-null	int64
4	Department	1470 non-null	object
5	DistanceFromHome	1470 non-null	int64
6	Education	1470 non-null	int64
7	EducationField	1470 non-null	object
8	EnvironmentSatisfaction	1470 non-null	int64
9	Gender	1470 non-null	object

10	HourlyRate	1470	non-null	int64
11	JobInvolvement	1470	non-null	int64
12	JobLevel	1470	non-null	int64
13	JobRole	1470	non-null	object
14	JobSatisfaction	1470	non-null	int64
15	MaritalStatus	1470	non-null	object
16	MonthlyIncome	1470	non-null	int64
17	MonthlyRate	1470	non-null	int64
18	NumCompaniesWorked	1470	non-null	int64
19	OverTime	1470	non-null	object
20	PercentSalaryHike	1470	non-null	int64
21	PerformanceRating	1470	non-null	int64
22	RelationshipSatisfaction	1470	non-null	int64
23	StockOptionLevel	1470	non-null	int64
24	TotalWorkingYears	1470	non-null	int64
25	TrainingTimesLastYear	1470	non-null	int64
26	WorkLifeBalance	1470	non-null	int64
27	YearsAtCompany	1470	non-null	int64
28	YearsInCurrentRole	1470	non-null	int64
29	YearsSinceLastPromotion	1470	non-null	int64
30	YearsWithCurrManager	1470	non-null	int64

dtypes: int64(23), object(8)
memory usage: 356.1+ KB

with above details,we can observe all the datatypes of each column along with memory consumption of dataset

#to get statistics information
df.describe()

	Age	DailyRate	...	YearsSinceLastPromotion
YearsWithCurrManager				
count	1470.000000	1470.000000	...	1470.000000
mean	36.923810	802.485714	...	2.187755
std	9.135373	403.509100	...	3.222430
min	18.000000	102.000000	...	0.000000
25%	30.000000	465.000000	...	0.000000
50%	36.000000	802.000000	...	1.000000
75%	43.000000	1157.000000	...	3.000000
max	60.000000	1499.000000	...	15.000000

[8 rows x 23 columns]

with above information,we can see there no missing values in any column and we can also observe the fact that many columns contains a minimum value of even 0 value,which means the value can even go with zero amount as well.

```
#to check all the values count in dependent variable 'wins(W)'  
df['Attrition'].value_counts()
```

```
No      1233  
Yes      237  
Name: Attrition, dtype: int64
```

with this,we can say that its a imbalanced data

```
#to check whether any null values present in dataframe  
df.isnull().sum()
```

```
Age      0  
Attrition      0  
BusinessTravel      0  
DailyRate      0  
Department      0  
DistanceFromHome      0  
Education      0  
EducationField      0  
EnvironmentSatisfaction      0  
Gender      0  
HourlyRate      0  
JobInvolvement      0  
JobLevel      0  
JobRole      0  
JobSatisfaction      0  
MaritalStatus      0  
MonthlyIncome      0  
MonthlyRate      0  
NumCompaniesWorked      0  
OverTime      0  
PercentSalaryHike      0  
PerformanceRating      0  
RelationshipSatisfaction      0  
StockOptionLevel      0  
TotalWorkingYears      0  
TrainingTimesLastYear      0  
WorkLifeBalance      0  
YearsAtCompany      0  
YearsInCurrentRole      0  
YearsSinceLastPromotion      0  
YearsWithCurrManager      0  
dtype: int64
```

So,there are no null values in the dataset

EDA

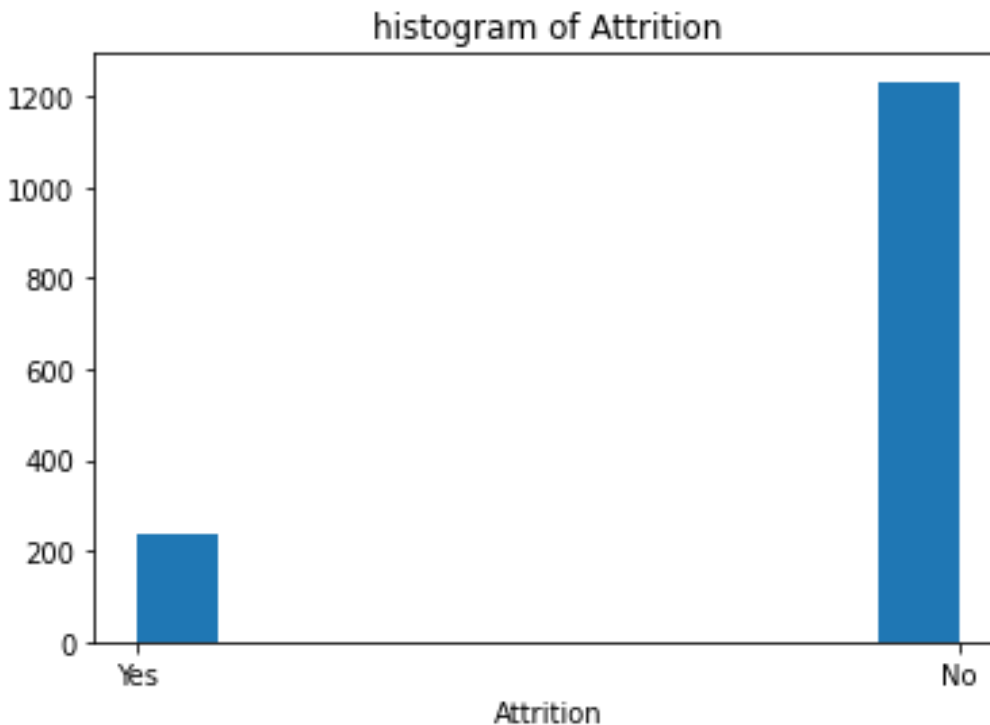
#visualize the histogram of wins

```
plt.hist(df['Attrition'])
```

```
plt.xlabel('Attrition')
```

```
plt.title('histogram of Attrition')
```

```
Text(0.5, 1.0, 'histogram of Attrition')
```




we can clearly see that "NO" values are more occurring than "YES" values. SO, we can deal with this kind of imbalanced data by using a F1-score as our metrics evaluation method after training the model

```
import warnings
```

```
warnings.filterwarnings('ignore') #to remove warning messages
```

```
sns.factorplot('Gender', data=df, kind="count") ## Let's check gender
```

```
<seaborn.axisgrid.FacetGrid at 0x7f21666e6390>
```




So, there are more records on male employee than female employee

Now let's separate the genders by MaritalStatus using 'hue' argument

```
sns.factorplot('Gender', data=df, kind="count", hue='MaritalStatus')
```

```
<seaborn.axisgrid.FacetGrid at 0x7f2165c875d0>
```




So, most records are on employees who are married and minimum records are on divorced employees.

Now let's separate the BusinessTravel by Department using 'hue' argument

```
sns.factorplot('BusinessTravel', data=df, kind="count", hue='Department')
```


```
<seaborn.axisgrid.FacetGrid at 0x7f216628bdd0>
```



So, employees are more who travels rarely and most employees are in Research & Development department


```
a = sns.factorplot('JobRole', data=df, kind="count")  
a.set_xticklabels(rotation=90)
```

```
<seaborn.axisgrid.FacetGrid at 0x7f2165b5f510>
```

So, most employees are working in as Sales Executive

```
a = sns.factorplot('EducationField', data=df, kind="count")  
a.set_xticklabels(rotation=60)
```



So,there are more employees whose education field has been the "life sciences"

#Since there are objective type columns,so we can use Label encoder method

```
from sklearn.preprocessing import LabelEncoder
```

```
df['Attrition'] = LabelEncoder().fit_transform(df['Attrition'])
df['BusinessTravel'] = LabelEncoder().fit_transform(df['BusinessTravel'])
df['Department'] = LabelEncoder().fit_transform(df['Department'])
df['EducationField'] = LabelEncoder().fit_transform(df['EducationField'])
df['JobRole'] = LabelEncoder().fit_transform(df['JobRole'])
df['MaritalStatus'] = LabelEncoder().fit_transform(df['MaritalStatus'])
df['OverTime'] = LabelEncoder().fit_transform(df['OverTime'])
df['Gender'] = LabelEncoder().fit_transform(df['Gender'])
```

```
df.head()
```

	Age	Attrition	...	YearsSinceLastPromotion	YearsWithCurrManager
0	41	1	...	0	5
1	49	0	...	1	7
2	37	1	...	0	0
3	33	0	...	3	0
4	27	0	...	2	2

[5 rows x 31 columns]

Correlation

df.corr()

	Age	...	YearsWithCurrManager
Age	1.000000	...	0.202089
Attrition	-0.159205	...	-0.156199
BusinessTravel	0.024751	...	-0.022636
DailyRate	0.010661	...	-0.026363
Department	-0.031882	...	0.034282
DistanceFromHome	-0.001686	...	0.014406
Education	0.208034	...	0.069065
EducationField	-0.040873	...	-0.004130
EnvironmentSatisfaction	0.010146	...	-0.004999
Gender	-0.036311	...	-0.030599
HourlyRate	0.024287	...	-0.020123
JobInvolvement	0.029820	...	0.025976
JobLevel	0.509604	...	0.375281
JobRole	-0.122427	...	-0.041150
JobSatisfaction	-0.004892	...	-0.027656
MaritalStatus	-0.095029	...	-0.038570
MonthlyIncome	0.497855	...	0.344079
MonthlyRate	0.028051	...	-0.036746
NumCompaniesWorked	0.299635	...	-0.110319
Overtime	0.028062	...	-0.041586
PercentSalaryHike	0.003634	...	-0.011985
PerformanceRating	0.001904	...	0.022827
RelationshipSatisfaction	0.053535	...	-0.000867
StockOptionLevel	0.037510	...	0.024698
TotalWorkingYears	0.680381	...	0.459188
TrainingTimesLastYear	-0.019621	...	-0.004096
WorkLifeBalance	-0.021490	...	0.002759
YearsAtCompany	0.311309	...	0.769212
YearsInCurrentRole	0.212901	...	0.714365
YearsSinceLastPromotion	0.216513	...	0.510224
YearsWithCurrManager	0.202089	...	1.000000

[31 rows x 31 columns]

#with target variable

df.corr()['Attrition'].sort_values()

TotalWorkingYears	-0.171063
JobLevel	-0.169105
YearsInCurrentRole	-0.160545
MonthlyIncome	-0.159840
Age	-0.159205
YearsWithCurrManager	-0.156199
StockOptionLevel	-0.137145
YearsAtCompany	-0.134392
JobInvolvement	-0.130016
JobSatisfaction	-0.103481
EnvironmentSatisfaction	-0.103369
WorkLifeBalance	-0.063939
TrainingTimesLastYear	-0.059478
DailyRate	-0.056652
RelationshipSatisfaction	-0.045872
YearsSinceLastPromotion	-0.033019
Education	-0.031373
PercentSalaryHike	-0.013478
HourlyRate	-0.006846
BusinessTravel	0.000074
PerformanceRating	0.002889
MonthlyRate	0.015170
EducationField	0.026846
Gender	0.029453
NumCompaniesWorked	0.043494
Department	0.063991
JobRole	0.067151
DistanceFromHome	0.077924
MaritalStatus	0.162070
OverTime	0.246118
Attrition	1.000000

Name: Attrition, dtype: float64

looks like there is less correlation between target variables and all independent variables

#plotting the correlation matrix using heatmap

```
corr_matrix = df.corr()
plt.figure(figsize=[20,13])
sns.heatmap(corr_matrix,annot=True)
plt.title("Correlation matrix")
plt.show()
```




```

df3 = df.iloc[:,16:24]
df4 = df.iloc[:,24:]

column_list = df1.columns.values
n_col = 31
n_row = 25
plt.figure(figsize = (2*n_col,4*n_col))
for i in range(len(column_list)):
    plt.subplot(n_row,n_col,i+1)
    sns.boxplot(df[column_list[i]],color='orange',orient='h')
    plt.tight_layout()

```




No outliers present in above columns

```

column_list = df2.columns.values
n_col = 31
n_row = 25
plt.figure(figsize = (2*n_col,4*n_col))
for i in range(len(column_list)):
    plt.subplot(n_row,n_col,i+1)
    sns.boxplot(df2[column_list[i]],color='orange',orient='h')
    plt.tight_layout()

```



no Outliers present in above graph

```


column_list = df3.columns.values
n_col = 31
n_row = 25
plt.figure(figsize = (2*n_col,4*n_col))

```

```

for i in range(len(column_list)):
    plt.subplot(n_row,n_col,i+1)
    sns.boxplot(df3[column_list[i]],color='orange',orient='h')
    plt.tight_layout()

```




Now, we can see there are outliers in "monthlyincome" along with few outliers appearing in other three columns.

```

column_list = df4.columns.values
n_col = 31
n_row = 25
plt.figure(figsize = (2*n_col,4*n_col))
for i in range(len(column_list)):
    plt.subplot(n_row,n_col,i+1)
    sns.boxplot(df4[column_list[i]],color='orange',orient='h')
    plt.tight_layout()

```



with above observation, we can see that "TotalWorkingyears", "YearsAtcompany", "Yearssincelastpromotion" columns have more number of outliers. whereas columns like "trainingTmeslastyear", "YearsinCurrentRole", "Yearswthcurrmanager" also contains a few outliers in it.

checking with Skewness

```
df.skew().sort_values()
```

BusinessTravel	-1.439006
WorkLifeBalance	-0.552480
JobInvolvement	-0.498419
Gender	-0.408665
JobRole	-0.357270
JobSatisfaction	-0.329672
EnvironmentSatisfaction	-0.321654
RelationshipSatisfaction	-0.302828
Education	-0.289681
MaritalStatus	-0.152175
HourlyRate	-0.032311
DailyRate	-0.003519
MonthlyRate	0.018578
Department	0.172231
Age	0.413286
EducationField	0.550371
TrainingTimesLastYear	0.553124
PercentSalaryHike	0.821128
YearsWithCurrManager	0.833451
YearsInCurrentRole	0.917363
DistanceFromHome	0.958118
OverTime	0.964489
StockOptionLevel	0.968980
JobLevel	1.025401
NumCompaniesWorked	1.026471
TotalWorkingYears	1.117172
MonthlyIncome	1.369817
YearsAtCompany	1.764529
Attrition	1.844366
PerformanceRating	1.921883
YearsSinceLastPromotion	1.984290

dtype: float64

by setting up the threshold for skewness in the range of -1 to +1

Data cleaning

With the above observations,

1) Columns like "PerformanceRating" and "BusinessTravel" have less correlation with target variable and also have skewness of more than 1.

2) Columns like "YearsSinceLastPromotion", "MonthlyIncome", "YearsAtCompany" and "TotalWorkingYears" have more number of outliers and also have skewness of more than 1.

So, we can remove all these columns from dataframe.

`df.columns`


```
Index(['Age', 'Attrition', 'BusinessTravel', 'DailyRate', 'Department',
      'DistanceFromHome', 'Education', 'EducationField',
      'EnvironmentSatisfaction', 'Gender', 'HourlyRate', 'JobInvolvement',
      'JobLevel', 'JobRole', 'JobSatisfaction', 'MaritalStatus',
      'MonthlyIncome', 'MonthlyRate', 'NumCompaniesWorked', 'OverTime',
      'PercentSalaryHike', 'PerformanceRating', 'RelationshipSatisfaction',
      'StockOptionLevel', 'TotalWorkingYears', 'TrainingTimesLastYear',
      'WorkLifeBalance', 'YearsAtCompany', 'YearsInCurrentRole',
      'YearsSinceLastPromotion', 'YearsWithCurrManager'],
      dtype='object')

df =
df.drop(['BusinessTravel', 'PerformanceRating', 'YearsSinceLastPromotion', 'Mont
hlyIncome', 'YearsAtCompany', 'TotalWorkingYears'], axis=1)

df.head()
```

	Age	Attrition	...	YearsInCurrentRole	YearsWithCurrManager
0	41	1	...	4	5
1	49	0	...	7	7
2	37	1	...	0	0
3	33	0	...	7	0
4	27	0	...	2	2

[5 rows x 25 columns]

Removing the outliers

#with z-score technique

```
from scipy.stats import zscore
import numpy as np
z_n = np.abs(zscore(df))
z_n.shape

(1470, 25)

threshold = 3

df_new = df[(z_n < threshold).all(axis=1)]
print(df.shape)
print(df_new.shape)

#printing the total number of dropped rows
print(df.shape[0]-df_new.shape[0])

(1470, 25)
(1445, 25)
25
```

calculating the percentage of data loss

```
loss_perc = (df.shape[0]-df_new.shape[0])/df.shape[0]*100
print("data loss percentage:{:.2f}%".format(loss_perc))
```

data loss percentage:1.70%

there is a 1.70% data loss which is good to proceed for further

Dividing the data into features and vectors

```
y = df_new[["Attrition"]]
x = df_new.drop(["Attrition"],axis=1)
```

Checking with multicollinearity

```
from statsmodels.stats.outliers_influence import variance_inflation_factor
def calc_vif_value(X):
```

```
    # Calculating VIF
```

```
    vif = pd.DataFrame()
```

```
    vif["variables"] = X.columns
```

```
    vif["VIF"] = [variance_inflation_factor(X.values, i) for i in
range(X.shape[1])]

    return(vif)
```

```
calc_vif_value(x)
```

	variables	VIF
0	Age	24.355684
1	DailyRate	4.965145
2	Department	12.651258
3	DistanceFromHome	2.303707
4	Education	9.297401
5	EducationField	3.824860
6	EnvironmentSatisfaction	7.002248
7	Gender	2.500496
8	HourlyRate	10.835432
9	JobInvolvement	14.317297
10	JobLevel	7.207210
11	JobRole	8.065798
12	JobSatisfaction	6.911987
13	MaritalStatus	5.642976
14	MonthlyRate	5.011830
15	NumCompaniesWorked	2.527821
16	OverTime	1.432536
17	PercentSalaryHike	15.994998
18	RelationshipSatisfaction	7.110562
19	StockOptionLevel	3.259987
20	TrainingTimesLastYear	5.603847
21	WorkLifeBalance	14.713868

```

22         YearsInCurrentRole    5.285521
23         YearsWithCurrManager   5.214726

```

So obviously we do have multicollinearity exist

Reducing the multicollinearity

```

x["Age_PercentSalaryHike"] = df.apply(lambda x_1: x_1['Age'] -
x_1['PercentSalaryHike'],axis=1)
x["WorkLifeBalance_JobInvolvement"] = df.apply(lambda x_1:
x_1['WorkLifeBalance'] - x_1['JobInvolvement'],axis=1)
x["Department_HourlyRate"] = df.apply(lambda x_1: x_1['Department'] -
x_1['HourlyRate'],axis=1)

x= x.drop(["Age", "PercentSalaryHike"],axis=1)
x = x.drop(["WorkLifeBalance", "JobInvolvement"],axis=1)
x = x.drop(["Department", "HourlyRate"],axis=1)

```

```
calc_vif_value(x)
```

	variables	VIF
0	DailyRate	4.834241
1	DistanceFromHome	2.283203
2	Education	8.764730
3	EducationField	3.748290
4	EnvironmentSatisfaction	6.757479
5	Gender	2.460481
6	JobLevel	6.678204
7	JobRole	4.130981
8	JobSatisfaction	6.635422
9	MaritalStatus	5.376706
10	MonthlyRate	4.890322
11	NumCompaniesWorked	2.487814
12	OverTime	1.429433
13	RelationshipSatisfaction	6.831798
14	StockOptionLevel	3.143458
15	TrainingTimesLastYear	5.393178
16	YearsInCurrentRole	5.268614
17	YearsWithCurrManager	5.205503
18	Age_PercentSalaryHike	8.369497
19	WorkLifeBalance_JobInvolvement	1.018102
20	Department_HourlyRate	9.624048

Now,its much better with the previous values

Transforming data to eliminate skewness

```

from sklearn.preprocessing import power_transform
x = power_transform(x)
x[:5]

```

```

array([[ 0.75507026, -1.49244581, -0.90978241, -0.96098587, -0.71378927,
        -1.22828358,  0.26777544,  1.03459175,  1.19013098,  1.23315508,
         0.74457901,  1.62050975,  1.58075594, -1.53026323, -1.08557575,
        -2.58064427,  0.23121958,  0.52529484,  0.86721048, -1.98951411,
        -1.32070858],
       [-1.34217869,  0.24626587, -1.76813882, -0.96098587,  0.20372452,
         0.81414424,  0.26777544,  0.61686941, -0.71519503, -0.13583322,
         1.4009283 , -0.56555193, -0.63260872,  1.22733082,  0.54943443,
         0.21803647,  0.93302132,  0.95965359,  0.47905416,  0.96589368,
         0.19850969],
       [ 1.33978769, -1.02912668, -0.90978241,  1.2458677 ,  1.21074448,
         0.81414424, -1.16284055, -1.00752372,  0.19231026,  1.23315508,
        -1.89355609,  1.27209045,  1.58075594, -0.70356727, -1.08557575,
         0.21803647, -1.5857131 , -1.54026201,  0.08168088,  0.96589368,
        -1.27488276],
       [ 1.37963762, -0.69799381,  1.0867191 , -0.96098587,  1.21074448,
        -1.22828358, -1.16284055,  0.61686941,  0.19231026, -0.13583322,
         1.1943781 , -0.56555193,  1.58075594,  0.22149186, -1.08557575,
         0.21803647,  0.93302132, -1.54026201,  0.08168088, -0.03479813,
         0.44782839],
       [-0.46245115, -1.02912668, -1.76813882,  0.69439463, -1.52351384,
         0.81414424, -1.16284055, -1.00752372, -0.71519503, -0.13583322,
         0.38010263,  1.76464883, -0.63260872,  1.22733082,  0.54943443,
         0.21803647, -0.43499851, -0.39322447, -0.64210101, -0.03479813,
         1.27603937]])

```

```

x_n1 = pd.DataFrame(x)
x_n1.skew()

```

```

0    -0.197102
1    -0.007307
2    -0.099535
3    -0.009425
4    -0.206012
5    -0.414570
6     0.106123
7    -0.323834
8    -0.213674
9    -0.157826
10   -0.183638
11     0.015440
12     0.949133
13   -0.195154
14     0.089303
15     0.057609
16   -0.071192
17   -0.077428
18     0.037999
19     0.000924

```

```
20      0.109972
dtype: float64
```

Its much better with skewness values compared to previous values

Selecting a best random state

#importing a necessary libraries

```
from sklearn.linear_model import LogisticRegression
from sklearn.metrics import accuracy_score
from sklearn.metrics import confusion_matrix,classification_report
from sklearn.model_selection import train_test_split
```

```
accu_score = 0
maxRS = 0 #best random state value for max accuracy_score
for i in range(1,200):
    x_train,x_test,y_train,y_test = train_test_split(x,y,test_size
=.20,random_state = i)
    LR = LogisticRegression()
    LR.fit(x_train,y_train)
    predic = LR.predict(x_test)
    acc = accuracy_score(y_test,predic)
    if acc > accu_score:
        accu_score = acc
        maxRS = i
print("Best accuracy_score is ",accu_score," on Random state ",maxRS)
```

```
Best accuracy_score is  0.9100346020761245  on Random state  71
```

```
x_train,x_test,y_train,y_test = train_test_split(x,y,test_size
=.20,random_state = 71)
```

```
print(x_train.shape,y_train.shape)
```

```
(1156, 21) (1156, 1)
```

```
print(x_test.shape,y_test.shape)
```

```
(289, 21) (289, 1)
```

To predict the Attrition status,we can consider Logistic regressiom model

```
LR = LogisticRegression()
LR.fit(x_train,y_train)
predic_1 = LR.predict(x_test)
print("accuracy_score is {:.3f}".format(accuracy_score(y_test,predic_1)))
```

#Since the class imbalance exists so, we can use F1-score as evaluation metrics

```
from sklearn.metrics import f1_score
print("f1_score is {:.3f}".format(f1_score(y_test,predic_1)))
```

#confusion matrix and classification report

```
print(confusion_matrix(y_test,predic_1))
print(classification_report(y_test,predic_1))
```

accuracy_score is 0.910

f1_score is 0.458

```
[[252  1]
```

```
 [ 25 11]]
```

	precision	recall	f1-score	support
0	0.91	1.00	0.95	253
1	0.92	0.31	0.46	36
accuracy			0.91	289
macro avg	0.91	0.65	0.70	289
weighted avg	0.91	0.91	0.89	289

#using DecisionTree classifier

```
from sklearn.tree import DecisionTreeClassifier
```

```
DT = DecisionTreeClassifier()
```

```
DT.fit(x_train,y_train)
```

```
predic_2 = DT.predict(x_test)
```

```
print("accuracy_score is {:.3f}".format(accuracy_score(y_test,predic_2)))
```

#Since the class imbalance exists so, we can use F1-score as evaluation metrics

```
from sklearn.metrics import f1_score
```

```
print("f1_score is {:.3f}".format(f1_score(y_test,predic_2)))
```

#confusion matrix and classification report

```
print(confusion_matrix(y_test,predic_2))
```

```
print(classification_report(y_test,predic_2))
```

accuracy_score is 0.806

f1_score is 0.349

```
[[218  35]
```

```
 [ 21 15]]
```

	precision	recall	f1-score	support
0	0.91	0.86	0.89	253
1	0.30	0.42	0.35	36
accuracy			0.81	289
macro avg	0.61	0.64	0.62	289
weighted avg	0.84	0.81	0.82	289

#Using KNN Classifier

```
from sklearn.neighbors import KNeighborsClassifier
```

```
KN = KNeighborsClassifier()  
KN.fit(x_train,y_train)  
predic_3 = KN.predict(x_test)  
print("accuracy_score is {:.3f}".format(accuracy_score(y_test,predic_3)))
```

#Since the class imbalance exists so, we can use F1-score as evaluation metrics

```
from sklearn.metrics import f1_score  
print("f1_score is {:.3f}".format(f1_score(y_test,predic_3)))
```

#confusion matrix and classification report

```
print(confusion_matrix(y_test,predic_3))  
print(classification_report(y_test,predic_3))
```

accuracy_score is 0.879

f1_score is 0.286

```
[[247   6]  
 [ 29   7]]
```

	precision	recall	f1-score	support
0	0.89	0.98	0.93	253
1	0.54	0.19	0.29	36
accuracy			0.88	289
macro avg	0.72	0.59	0.61	289
weighted avg	0.85	0.88	0.85	289

#Using Random forest Classifier

```
from sklearn.ensemble import RandomForestClassifier
```

```
RF = RandomForestClassifier()  
RF.fit(x_train,y_train)  
predic_4 = RF.predict(x_test)  
print("accuracy_score is {:.3f}".format(accuracy_score(y_test,predic_4)))
```

#Since the class imbalance exists so, we can use F1-score as evaluation metrics

```
from sklearn.metrics import f1_score  
print("f1_score is {:.3f}".format(f1_score(y_test,predic_4)))
```

#confusion matrix and classification report

```
print(confusion_matrix(y_test,predic_4))  
print(classification_report(y_test,predic_4))
```

accuracy_score is 0.896

f1_score is 0.318

```
[[252  1]
 [ 29  7]]
```

	precision	recall	f1-score	support
0	0.90	1.00	0.94	253
1	0.88	0.19	0.32	36
accuracy			0.90	289
macro avg	0.89	0.60	0.63	289
weighted avg	0.89	0.90	0.87	289

f1_score and accuracy is high for logistic regression model

Cross Validation to check with overfitting

```
from sklearn.model_selection import cross_val_score
```

```
#for Logistic regression model
```

```
L_cr = cross_val_score(LR,x,y,cv=5)
```

```
print("cross validation for Logistics regression model : ",L_cr.mean())
```

```
cross validation for Logistics regression model : 0.8622837370242215
```

```
L_cr = cross_val_score(KN,x,y,cv=5)
```

```
print("cross validation for KNN classifier model : ",L_cr.mean())
```

```
cross validation for KNN classifier model : 0.8401384083044983
```

```
L_cr = cross_val_score(DT,x,y,cv=5)
```

```
print("cross validation for Decision tree model : ",L_cr.mean())
```

```
cross validation for Decision tree model : 0.7647058823529411
```

```
L_cr = cross_val_score(RF,x,y,cv=5)
```

```
print("cross validation for Random forest model : ",L_cr.mean())
```

```
cross validation for Random forest model : 0.853287197231834
```

here also,the CV score is more for Logistic regression model. So, I will take my final model as logistic regression model

Hyper parameter tuning

```
from sklearn.model_selection import GridSearchCV
```

```
params_list = {'penalty':['l1','l2','elasticnet','none'],
               'solver':['newton-cg','lbfgs','liblinear','sag','saga'],
               'multi_class':['auto','ovr','multinomial']}
```



```
GCV = GridSearchCV(LogisticRegression(),params_list,cv=5,scoring="accuracy")
GCV.fit(x_train,y_train)
```

```
#printing the best parameter
```

```
GCV.best_params_
```

```
{'multi_class': 'auto', 'penalty': 'l2', 'solver': 'newton-cg'}
```

```
#now predicting with best parameters
```

```
GCV_predic = GCV.best_estimator_.predict(x_test)
```

```
#getting the final accuracy
```

```
accuracy_score(y_test,GCV_predic)
```

```
#f1_score(y_test,GCV_predic)
```

```
0.9100346020761245
```


it was almost same accuracy and f1_score compared with before tuning the model

```
from sklearn.metrics import plot_roc_curve
```

```
plot_roc_curve(GCV.best_estimator_,x_test,y_test)
```

```
plt.title("ROC AUC plot")
```

```
plt.show()
```



from the above graph we got AUC score is 90%.

```
import joblib
joblib.dump(GCV.best_estimator_, "HR_Analytics_Project.pkl")
['HR_Analytics_Project.pkl']
```